R1CS Programming
ZK0x04 Workshop Notes

Daniel Lubarov   Brendan Farmer

October 24, 2019

Contents

1 Multiplicative inverse

Deterministically computing $1/x$ in an R1CS circuit would be expensive. Instead, we can have the prover compute $1/x$ outside of the circuit and supply the result as a witness element, which we will call x_{inv}. To verify the result, we enforce

$$(x)(x_{\text{inv}}) = (1) \tag{1}$$

2 Zero testing

To assert $x = 0$, we simply enforce

$$(x)(1) = (0) \tag{2}$$

Asserting $x \neq 0$ is similarly easy: we compute $1/x$ (non-deterministically, as in Section 1). The result can be ignored; the mere fact that an inverse exists implies $x \neq 0$.

On the other hand, if we want to *evaluate*

$$y := \begin{cases} 0 & \text{if } x = 0, \\ 1 & \text{otherwise,} \end{cases} \tag{3}$$

we can do so by introducing another variable, m, and enforcing

$$(x)(m) = (y), \tag{4}$$
$$(1 - y)(x) = (0). \tag{5}$$

Outside of the circuit, the prover generates y as in Equation 3, and generates m as

$$m := \begin{cases} 1 & \text{if } x = 0, \\ y/x & \text{otherwise.} \end{cases} \tag{6}$$

This method is from [1].

We can also use this technique to test for equality, since $x = y$ if and only if $x - y = 0$.

3 Binary

To assert $b \in \{0, 1\}$, we enforce

$$(b)(b - 1) = (0). \tag{7}$$

To "split" a field element x into its binary encoding, $(b_1, \ldots, b_n)$, we have the prover generate the binary encoding out-of-band. We then verify it by applying Equation 7 to each b_i, and enforcing

$$(x)(1) = \left(\sum_{i=0}^{n-1} 2^i b_i \right), \tag{8}$$

assuming a little-ending ordering of the bits.

Note that Equation 8 permits two encodings of certain field elements. In $\mathbb{F}_{13}$ for example, the element 1 can be represented as either 0001 or 1110. If a canonical encoding is required, we can prevent "overflowing" encodings by asserting that $(b_1, \ldots, b_n) < |F|$. Such binary comparisons are covered in Section 10.

To "join" a sequence of bits into the field element they encode, we simply take a weighted sum of the bits:

$$x := \sum_{i=0}^{n-1} 2^i b_i. \tag{9}$$

This does not require any constraints, unless we don't know whether $(b_1, \ldots, b_n)$ is the canonical encoding of some field element, in which case we may want a comparison to assert that overflow does not occur.

4 Selection

Suppose we have a boolean value s, and we wish to compute

$$z := \begin{cases} x & \text{if } s = 0, \\ y & \text{if } s = 1. \end{cases} \tag{10}$$

We can compute this as

$$z := x + s(y - x). \tag{11}$$

This requires two constraints: one "is boolean" assertion (Equation 7), assuming s was not already known to be boolean, and another for the multiplication.

5 Random access

Suppose we wish to access the ith item of some sequence $(x_0, \ldots, x_n)$. There are several variations of this problem, but for the sake of brevity, we will assume that

1. The sequence length is a power of two: $n = 2^k$.

2. The items are individual field elements (rather than, say, elliptic curve points).

3. Each item x_j is a constant.

We start by splitting i into its binary encoding, $(i_0, \ldots, i_k)$. One may think of $(i_0, \ldots, i_k)$ as a path through a binary tree, where i_j indicates the direction of descent at height j. For example, say $n = 8$ (thus $k = 3$) and $i = 5 = 101$. This corresponds to the following binary tree:

Now, we will describe a few methods of random access which build upon this mental model.

5.1 Repeated selection method

First, we could perform a selection operation (Section 4) for each non-leaf node. In the example, we would perform the following operations:
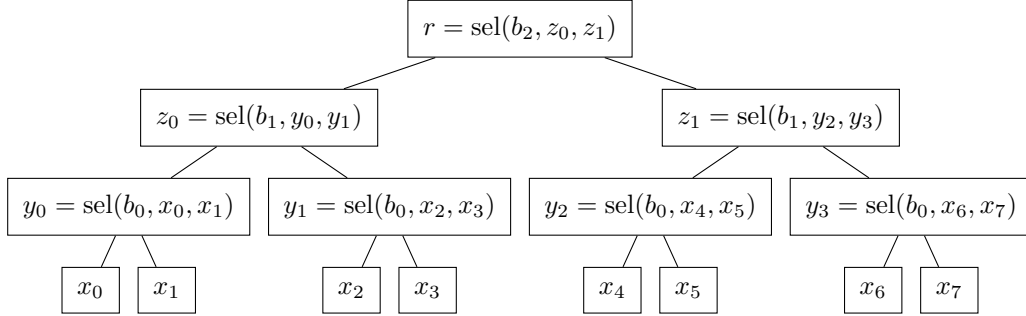
after which the root node r will hold x_i. This involves $n-1$ (or $2^k - 1$) selection operations, since a perfect tree with n leaves contains $n-1$ non-leaf nodes. However, if $(x_0, \ldots, x_n)$ are constant then selecting between two leaves becomes a free linear combination, which brings the cost down to $n/2 - 1$ (or $2^{k-1} - 1$) constraints.

5.2 Sum-of-conditions method

Alternatively, we could construct an expression for each node which encodes the conditions under which that node, or one of its ancestors, is selected. In the example, we have

Then to access x_i, we take the sum of each leaf multiplied by its selection condition. Note that this final multiplication is free, given our assumption that $(x_0, \ldots, x_n)$ are constants.

Computing these conditions naively would require $n - 1$ constraints, since we would perform one multiplication per non-root node. We can achieve much greater efficiency, however, by distributing the products and reusing intermediate products. In the example, we would compute just four products: $b_0 b_1$, $b_0 b_2$, $b_1 b_2$, and $b_0(b_1 b_2)$. In general, the cost of this approach is $2^k - k - 1$ constraints.

5.3 Hybrid method

Unfortunately, both of the aforementioned methods cost $\mathcal{O}(n)$ constraints. It turns out that we can do better by combining them. Pick some l, m such that $l + m = k$. We divide our conceptual tree into subtrees of size 2^l, and apply the sum-of-conditions method to each subtree. This costs $2^l - l - 1$ constraints independent of the number of subtrees, since each subtree uses the exact same set of conditions.

Next, we invoke the selection method to select one of the 2^m subtree results. This costs another $2^m - 1$ constraints, for a total cost of $2^m + 2^l - l - 2$. By picking $l = m = k/2$, we obtain a total cost of $\mathcal{O}(\sqrt{n})$ constraints.

6 Embedded curve operations

Embedded curves have several uses in SNARKs. A few examples are Schnorr signatures, Pedersen hashes, and recursive SNARK verifiers. Here we will focus on twisted Edwards curves such as Jubjub.

6.1 Addition

Recall the addition law for twisted Edwards curves,

$$(x_1, y_1) + (x_2, y_2) = \left(\frac{x_1 y_1 + y_1 x_2}{1 + d x_1 x_2 y_1 y_2}, \frac{y_1 y_2 - a x_1 x_2}{1 - d x_1 x_2 y_1 y_2} \right). \tag{12}$$

Applying the law directly takes 7 constraints: 4 for the products in the numerators, one for the product in the denominator, and one[1] for each of the two quotients. This can be optimized to 6 constraints as described in [2].

However, the operation becomes much cheaper when of the points is constant. Without loss of generality, suppose (x_1, y_1) is constant while (x_2, y_2) is variable. Then the numerators become "free" linear combinations, while the denominators require a single multiplication. The quotients add 2 constraints as before, resulting in 3 constraints total.

Finally, when doubling a point, the addition law simplifies to

$$[2](x, y) = \left(\frac{2xy}{ax^2 + y^2}, \frac{y^2 - ax^2}{2 - ax^2 - y^2} \right) \tag{13}$$

which requires 5 constraints.

6.2 Multiplication

Consider the problem of computing $[s]p$, where p is a constant point and s is a variable scalar. A simple, reasonably efficient approach is known as multiplication by doubling. We split s into its binary encoding, $(s_0, \ldots, s_n)$, then compute

$$[s]p = \sum_{i=0}^{n-1} [s_i][2^i]p. \tag{14}$$

Since p is constant, $[2^i]p$ can be precomputed. Thus for each i, we perform a conditional addition: if $s_i = 1$, we add $[2^i]p$ to an accumulator. Since adding a constant point is cheaper than adding a variable point, it is best to structure the computation as

```
sum ← 0;
for i ← 0 to n do
    q ← sum + [2^i]p;                    /* addition with a constant value */
    sum ← select(s_i, sum, q);
end
return sum;
```

[1] In general, computing a quotient $q \coloneqq x/y$ takes two constraints: $(y)(y_{\mathrm{inv}}) = (1)$ and $(x)(y_{\mathrm{inv}}) = (q)$. In this case, however, we can multiply both sides by the denominator since we know it will never be zero. This yields a single constraint: $(q)(y) = (x)$.

Ignoring the cost of splitting, this requires 5 constraints per bit: 3 for the constant addition, and 2 for selecting a 2-tuple.

6.2.1 Windowed multiplication

We can generalize Equation 14 by decomposing s into its base-2^w encoding, rather than base 2 specifically. Let $(s_0, \ldots, s_{n/w})$ be the base-2^w decomposition of s. Note that s_i can be expressed as a linear combination of bits of s, so the cost of decomposition does not change. We can rewrite $[s]p$ as

$$[s]p = \sum_{i=0}^{n/w-1} [s_i][(2^w)^i]p. \tag{15}$$

As before, $[(2^w)^i]p$ can be precomputed for each i. We can go a step further by precomputing $[s_i][(2^w)^i]p$ for each of 2^w possible values of s_i, then looking up the relevant value using the method described in Section 5.3.

In practice, a window size of 3 seems optimal. If we configure our random access gadget with $l = 2$ and $m = 1$, the lookup costs 3 constraints per 3-bit window.[2] The addition adds another 7 constraints per window, for a total cost of 10 constraints per window, or 3.33 constraints per bit.

7   2x2 switch

Suppose we wish to implement a switch with the following structure:

In particular, if $s = 0$ then the outputs should be identical the inputs: $(c, d) = (a, b)$. If $s = 1$ then the inputs should be swapped: $(c, d) = (b, a)$.

This requires two constraints: one "is boolean" assertion (Equation 7), and another for selecting the value of c (Equation 11). Once we have c, we can compute d as

$$d := a + b - c \tag{16}$$

which does not require any additional constraints.

8   Permutations

Say we want to verify that two sequences, $(x_1, \ldots, x_n)$ and $(y_1, \ldots, y_n)$, are permutations of one another. This can be done efficiently using routing networks, which used a fixed (for a fixed n) network of 2x2 switches.

AS-Waksman networks [3] are a particularly useful construction, since they support arbitrary permutation sizes. They use about $n \log_2(n) - n$ switches, which is close to the theoretical lower bound of $\log_2(n!)$.

---

[2]This is one more than the formula in Section 5.3 would imply, since the selection operation needs to be done for both coordinates of the curve point.

9   Sorting

Like permutation networks, sorting networks use a fixed network of gates. In particular, a sorting network is comprised of several 2x2 comparator gates, each of which takes two inputs and sorts them. It is theoretically possible to construct a sorting network for n elements using $\mathcal{O}(n \log n)$ gates [4], but practical constructions use $\mathcal{O}(n \log^2 n)$ gates. Since each comparison adds $\mathcal{O}(\log |F|)$ constraints, this approach is fairly expensive.

A better solution is to leverage non-determinism: instead of creating an R1CS circuit to sort a sequence, we have the prover supply the ordered sequence. Using a permutation network (Section 8), we can efficiently verify that the two sequences are permutations of one another. Then for each contiguous pair of elements in the ordered sequence, x_i and x_{i+1}, we assert $x_i \leq x_{i+1}$.

10   Comparisons

Suppose we want to evaluate $x < y$. For now, assume that both values fit in n bits, where $2^{n+1} < |F|$. We compute

$$z = 2^n + x - y. \tag{17}$$

Notice that 2^n has a single 1 bit with index n, which will be cleared if and only if $x < y$. Thus we split z into $n + 1$ bits, and the negation of the nth bit is our result.

If x and y are unbounded, then the aforementioned technique does not apply. A neat solution was described by Ahmed Kosba. First, we split x and y into their binary forms, $(x_0, \ldots, x_n)$ and $(y_0, \ldots, y_n)$. Next we divide each binary sequence into c-bit chunks, and join each chunk of bits into a field element (Equation 9).

The prover non-deterministically identifies the most significant chunk index at which x and y differ, if any. We add constraints to verify that the identified chunks do indeed differ, and that any more significant chunks are equal. Finally, we compare the two identified chunks using Equation 17.

If we choose $c, l \in \mathcal{O}(\sqrt{n})$, this approach requires $\mathcal{O}(\sqrt{n})$ constraints in addition to the cost of splitting the inputs (if they were not already in binary form). A couple optimizations are possible in particular circumstances:

1. To assert (not evaluate) $x < y$, we can split x non-canonically and split y canonically. The prover is forced to use x's canonical encoding anyway, otherwise $x_{\text{bin}} \geq |F| > y_{\text{bin}}$, making the assertion unsatisfiable.

2. To assert $x < c$ for some constant $c \ll |F|$, we can split x into just $\lceil \log_2 c \rceil$ bits.

References

[1] B. Parno, J. Howell, C. Gentry, and M. Raykova, "Pinocchio: Nearly practical verifiable computation," in *2013 IEEE Symposium on Security and Privacy*, pp. 238–252, IEEE, 2013.

[2] D. Hopwood, S. Bowe, T. Hornby, and N. Wilcox, "Zcash protocol specification,"

[3] B. Beauquier and E. Darrot, "On arbitrary size waksman networks and their vulnerability," *Parallel Processing Letters*, vol. 12, no. 03n04, pp. 287–296, 2002.

[4] M. Ajtai, J. Komlós, and E. Szemerédi, "An o(n log n) sorting network," in *Proceedings of the fifteenth annual ACM symposium on Theory of computing*, pp. 1–9, ACM, 1983.